

# DISFI (Distributed Infrastructure Stress and Fault Injector) — User Manual

## 1. Overview

DISFI (Distributed Infrastructure Stress and Fault Injector) is a configurable fault-injection framework designed to induce controlled stress conditions and resource exhaustion in distributed, stateful systems. The tool generates sustained background telemetry traffic while selectively injecting pathological workload patterns, including connection storms, read/write amplification, intentional connection leaks, and CPU saturation. DISFI is specifically designed to trigger internal failure modes such as thread pool saturation, file descriptor exhaustion, request timeouts, and transient unavailability. All injected faults and execution parameters are explicitly logged, enabling precise temporal alignment between injected stress conditions and observed system anomalies. The framework supports reproducible experimental campaigns for anomaly detection, predictive maintenance, and reliability analysis in distributed databases and streaming platforms.

It is intended for:

- Testing predictive anomaly detection
- Producing labeled anomalies for ML pipelines
- Stress-testing observability dashboards (Prometheus, Grafana)
- Validation of alerting rules

## 2. Command Line Usage

```
python traffic_gen.py \
--hosts 192.168.1.10,192.168.1.11 \
--rate 200 \
--threads 50 \
--mode read \
--duration-minutes 10 \
--log logs/file.log
```

### Mandatory parameters

Parameter	Meaning
--hosts	Comma-separated Cassandra IPs
--mode	read, write, mixed
--threads	Generator threads
--rate	Operations per second

### Optional parameters

Optional	Meaning
--duration-minutes	Auto-stop timeout
--storm	Enable Connection Storm scenario
--cpu-spike	Enable artificial CPU saturation

<b>Optional</b>	<b>Meaning</b>
--cluster-timeout	ms timeout for connection
--log	Persistent JSON logging

If duration-minutes is not provided, the script **runs indefinitely**.

## 3. Observability and Logging

The generator writes extended diagnostic info:

- Start timestamp
- Stop timestamp
- Effective duration
- All CLI parameters
- Number of operations executed
- Exception count
- Performance counters

Useful for:

- auditability
- experiment reproducibility
- integration with notebooks or dashboards

**Architecture and execution flow of the Cassandra stress traffic generator.**

